GradientCIO

INVESTMENT DUE DILIGENCE

Northstar Manager Comparison

Comparative performance and benchmark assessment

REVIEW

Manager comparison

PREPARED FOR

Investment Committee

DATA AS OF

2025-12-31

01 Executive conclusion

Evidence does not support a comparative conclusion

Key judgment — Coverage

Observation: Only one of four comparison inputs is available S1. Why it matters: Managers cannot be ranked on one input. Uncertainty: Missing return series may change the order. What would change the view: Return series for every candidate.

Key judgment — Benchmarks

Observation: Benchmark matches are suggested, not confirmed S1. Why it matters: Relative results depend on the right benchmark. Uncertainty: Mandate fit is unconfirmed. What would change the view: Confirmed benchmark mappings.

Key judgment — Reliability

Observation: Robustness across date windows was not tested S1. Why it matters: A single window can flatter one manager. Uncertainty: Results may depend on the start date. What would change the view: A date-window robustness run.

02 Data and benchmark coverage

Benchmark identity remains advisory for the reviewed candidates

Benchmark mapping	**Not available**	Canonical match unavailable S1

03 Risk and return

Risk and return comparison is unavailable

Not available. Governed comparison result missing S1.

Analysis — Comparison unavailable

Observation: Comparative metrics are unavailable S1. Why it matters: No ranking is supportable. Uncertainty: Missing evidence may differ by manager. What would change the view: A governed comparison result.

04 Relative performance and capture

Relative and capture evidence is unavailable

Not available. Relative-return result missing S1.

Analysis — Relative evidence missing

Observation: Relative evidence is unavailable S1. Why it matters: Benchmark-relative behavior is unknown. Uncertainty: Capture may vary by window. What would change the view: A governed relative-return result.

05 Correlation and robustness

Persistence and correlation evidence is unavailable

Not available. Governed robustness and correlation results missing S1.

Analysis — Robustness unavailable

Observation: Robustness evidence is unavailable S1. Why it matters: Persistence cannot be assessed. Uncertainty: Results may depend on dates. What would change the view: Governed window tests.

06 Follow-ups

One governed evidence request remains

01 **Return canonical benchmark and comparison evidence.**
Required to assess manager-relative performance S1.

A Appendix — Sources and methods

Saved illustrative sample evidence

Based on 1 of 4 evidence areas.

TAG	EVIDENCE	SOURCE	AS OF	CHECK
S1	Typed unavailable test response	get_gradient_capabilities	2025-12-31	Passed

Illustrative, Gradient Maintained — demo data, not the client's holdings or managers. All organizations, portfolios, managers and results in this public example are illustrative.

SOURCE AUDIT TRAIL

TAG	KEY PARAMETERS	DATA SCOPE	DIGEST
S1	illustrative sample	Illustrative sample data	sha256:aaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaa

Each digest is the SHA-256 fingerprint of the exact data the report used.

Governed evidence boundary

The report preserves the returned subject identity, benchmark mapping, observation window, units, methods, coverage status and request provenance S1. Missing performance, correlation, capture, simulation or regime fields remain unavailable. No local statistics, substitutions, rankings or neutral defaults are introduced.

Decision-use boundary

Typed unavailability limits the conclusion and drives the insufficient signal where most required slots are absent S1. The evidence supports follow-up collection only; operational diligence and committee decisions remain separate handoffs with their own governed evidence and validation contracts. Every unresolved field remains reader-visible in the delivered report.

Metric and digest interpretation

Historical claims must use returned realized metrics with their exact periods, while path-distribution statistics must remain explicitly simulated S1. Successful tool evidence records the SHA-256 digest of the saved response; a failed call instead uses the typed no-payload sentinel and preserves its error code and request ID in the validation field. These controls prevent simulated statistics or request identifiers from being presented as observed evidence. The source row also retains the tool name, key parameters, as-of date and data scope so a reviewer can reproduce the evidence boundary without treating an unavailable payload as a successful observation. Advisory benchmark selections remain visibly advisory throughout the report and handoff. This distinction tells reviewers which comparisons are server-matched and which still require independent benchmark confirmation.